

## CSSE2310 — 4.2

### Processes

# Abstraction

A process is:

- ▶ An instance of a program in execution.
  - ▶ There is one `/opt/local/bin/bash` on moss.
  - ▶ There could be many executing instances of bash
- ▶ An abstraction of a computer
  - ▶ Interactions (Eg with files) are via system calls to the kernel.
  - ▶ Other processes' resources are not visible.
  - ▶ Processes' memory is <sup>1</sup> not accessible by other processes
  - ▶ Other processes' cpu activity shouldn't influence your's.
  - ▶ Indirectly aware of other non-related processes
    - ▶ Eg Files on the filesystem.

---

<sup>1</sup>generally

# How?

Separating:

- ▶ Resources (eg open files)
  - ▶ Each process can have their own table of resources.
  - ▶ At lower level, files are identified by integers.
- ▶ Memory
  - ▶ Virtual memory (next week)
- ▶ CPU activity (From Week 4)
  - ▶ Whenever a cpu switches to kernel mode, registers are saved.
  - ▶ (possibly much later) when the process is put on the cpu, registers are restored.
  - ▶ The process is none the wiser.

# Process states

From the kernel point of view, a process could be in one of the following states:

- ▶ Running — The process is currently executing on the CPU
- ▶ Ready — The process could run but isn't (something else has the CPU).
- ▶ Blocked — Process is not ready because it is waiting for something.
  - ▶ eg: waiting to read from a file
  - ▶ sleep?
- ▶ Ended(?) — Process has exited or terminated and needs to be cleaned up
  - ▶ May have become a zombie

# The circle of life

Fork discussion

## fork()

fork()

- ▶ Asks the kernel to create a new (child) process
- ▶ Called in one process (the parent)
- ▶ Returns in two processes (two processes remember calling it).
- ▶ Parent gets the PID of the child.
  - ▶ or if no child was created (and sets errno).
- ▶ Child gets 0.

```
pid_t pid = fork();  
if (pid) {  
    // parent  
} else {  
    // child  
}
```


Memory aid: A parent can have many children (so it needs to know who is new). The child process only has one parent (and can getppid()).

## Memory starts as a copy

```
void X(void) {  
    int p=5;  
    if (fork()) {  
        ++p;  
    } else {  
        --p;  
    }  
    printf("%d\n", p);  
}  
  
int main(int argc, char** argv) {  
    X();  
    printf("Done");  
    return 0;  
}
```

See f1.c


Child processes are processes and can fork as well.  
See f2.c.






## Watch out for buffering

See f3.c vs f4.c.



## More care

The parent process could fork to make workers  
See f5.c



Hmmm

```
for (int i=0; i<20; ++i) {  
    if (!fork()) {  
        printf("Child\n");  
    }  
}
```

## fork bomb

The example on the previous slide will terminate eventually. Not always the case.

Difficult to stop:

- ▶ Limited number of simultaneous processes per user.
  - ▶ `ulimit -u` on bash
- ▶ Use `kill` or `pkill`?
  - ▶ In most cases<sup>2</sup> that needs to fork again.
  - ▶ Which you can't do
  - ▶ How do you know which process IDs to kill?<sup>3</sup>
- ▶ Can you do anything quickly enough?
  - ▶ Your useful programs will be competing with all those clones.

---

<sup>2</sup>Unless your shell has a built in `kill` (bash doesn't)

<sup>3</sup>`kill -9 -1?`

# Fork bomb consequences

Depends on the system configuration. On moss:

- ▶ May slow *your* processes down.
- ▶ *Probably* won't affect other users
- ▶ “But moss is being slow”
  - ▶ Moss almost always I/O bound.
  - ▶ Unless sysadmins say otherwise, Someone else fork bombing themselves is unlikely to be causing problems for you.

If you manage to do this to yourself.

And **it doesn't die down on its own**. Contact the helpdesk.

## Ending a process

`exit()` is a system call which ends the *current process*. ie you can't use it to end a different process.

Useful things:

- ▶ The parameter to `exit` is the “exit status” of the process.
- ▶ Any open output streams are flushed
- ▶ exit hooks are executed
  - ▶ `int atexit(void (*function)(void));`

Don't want those things? (eg if you think you have inherited unflushed buffers)

- ▶ `man _exit()`

Program crashes (or receives a signal<sup>4</sup>)

- ▶ `exit` doesn't happen and there is no exit status.

---

<sup>4</sup>Signals are later