

## CSSE2310/7231 — 5.2

### Processes (continued)

# ZOMBIES

See `zom1.c`.

# ZOMBIES

See `zom1.c`.

The system needs to keep a record of the what happened to a process in case the parent is interested.

- ▶ Did it exit normally?
  - ▶ with what status?
- ▶ It terminated rather than exiting normally.
  - ▶ Which signal caused that?

The memory and resources the process was using have already been released but part of the process still hangs around. A process in that state is called a **zombie**.

# Reaping a zombie

Reaping — when a process' parent has asked about COD.


- ▶ zombie will be removed

To reap, the parent process calls `wait()`. Wait blocks until either:

- ▶ A current child process becomes (or already was) a zombie.
- ▶ The parent has no child processes (returns error)

See `zom2.c`

See `zom2.c`



Zombies can't be killed. — There is no code left to handle a signal.

Zombies can't be killed. — There is no code left to handle a signal.

Q: How do you stop a zombie?



Zombies can't be killed. — There is no code left to handle a signal.

Q: How do you stop a zombie?

A: You drop a wait on it.

# Reaping and adoption

If a process

- ▶ is “alive”
- ▶ has zombie children
- ▶ doesn't reap them

then those zombies will stay on the system (the parent might ask about them eventually). For long running processes like servers, this could be a problem.

# Reaping and adoption

What if a process dies and has running child processes?

See `adopt.c`

# Reaping and adoption

What if a process dies and has running child processes?

See `adopt.c`

- ▶ Only the direct parent of a process can reap it.
- ▶ The children will continue to run.
- ▶ All that process' children will be adopted by process 1.

## Decoding “status”

The information returned by `wait` needs to be decoded.

- ▶ `WIFEXITED(status)` — true if the process exited normally
  - ▶ then `WEXITSTATUS(status)` — the exit status of the process.
- ▶ `WIFSIGNALED(status)` — true if the process was terminated by a signal.
  - ▶ `WTERMSIG(status)` — the signal which caused the process to terminate.

See `f7.c`

`Wait` can be used to check for things other than a process ending, but we don't need them now.

```
waitpid()
```

See `f8.c`

The third argument will be 0 for now.  
Later we may use `W_NOHANG`.

## Changing script



A process can change which program it is running.

```
int execl(char* path, char* arg0, char* arg1, char* arg2, ...)
```

- ▶ The last thing in the list must be a null pointer
- ▶ Replaces the old process image with a new one.
  - ▶ The old stack and heap are gone.
  - ▶ The old program instructions are gone.
  - ▶ Resources on the kernel side (eg files) are kept.
- ▶ returns `-1` if the operation fails
- ▶ never returns anything but `-1`.

See `execl.c`

# Pathing

- ▶ From `exec1.c`:  
`./a.out ls` doesn't work.
- ▶ To the contents of the `PATH` variable into account, use `exec1p`.
- ▶ In this course, you should always use the `p` forms of `exec`.





See `exec2.c`

```
int execvp(char* arg0, char** argv).
```

- ▶ More useful for varying numbers of commandline arguments.
- ▶ The last element in `argv`, must be a null pointer.

There is also `execvpe()`<sup>1</sup> but we don't need them in this course.

---

<sup>1</sup>and others

## Summary so far

So, ... in order to run another program ...

# Summary so far

So, ... in order to run another program ...

- ▶ we clone ourselves

# Summary so far

So, ... in order to run another program ...

- ▶ we clone ourselves
- ▶ replace the clone's memory with a new program.

# Summary so far

So, ... in order to run another program ...

- ▶ we clone ourselves
- ▶ replace the clone's memory with a new program.

Is this overly complicated or is there a reason for this?

# Summary so far

So, ... in order to run another program ...

- ▶ we clone ourselves
- ▶ replace the clone's memory with a new program.

Is this overly complicated or is there a reason for this?

Yes

# Signals

- ▶ Signals are very simple messages from the kernel to a process.
  - ▶ They don't carry information other than "One or more of some event has happened"
  - ▶ If a previous signal of that type hasn't been handled yet, additional signals<sup>2</sup> will be ignored.
- ▶ The kernel will send signals:
  - ▶ On its own initiative:
    - ▶ the process has accessed invalid memory — segfault (SIGSEGV)
    - ▶ the process tried to write to a destination that won't accept input (SIGPIPE)
    - ▶ ...
  - ▶ Because a process has asked it to:
    - ▶ shell `kill` command.
    - ▶ C `kill` system call.

---

<sup>2</sup>of that type

# kill?

- ▶ Processes have signal handling functions registered with the kernel.
  - ▶ When a signal is delivered, the kernel fakes a function call in the process (see “exceptions” last week).
- ▶ If your program doesn't set up new ones, the default handler will be used.
  - ▶ The default handler usually terminates the process.
- ▶ `kill` commands are actually “send a given signal to a process”.
  - ▶ ...often that results in the process not being alive.



## signal numbers

To see signals available on your system in bash:

```
kill -l (n.b. dash el)
```

```
1) SIGHUP 2) SIGINT 3) SIGQUIT 4) SIGILL 5) SIGTRAP  
6) SIGABRT 7) SIGBUS 8) SIGFPE 9) SIGKILL 10) SIGUSR1  
...
```

In code, use symbolic names:

eg SIGINT

# sigaction

From man sigaction

```
struct sigaction {  
    void      (*sa_handler)(int);  
    void      (*sa_sigaction)(int, siginfo_t *, void *);  
    sigset_t   sa_mask;  
    int       sa_flags;  
    void      (*sa_restorer)(void);  
};
```

There are older functions like:

- ▶ signal
- ▶ sigset

do not use them.

sig.c

See sig.c

See sig.c

How do I stop valgrind complaining about my sigaction variable?

- ▶ Add `memset(&sa, 0, sizeof(struct sigaction));` after declaration

See sig.c

How do I stop valgrind complaining about my sigaction variable?

- ▶ Add `memset(&sa, 0, sizeof(struct sigaction));` after declaration

SA\_RESTART?

- ▶ If a signal arrives while a system call is running, restart the sys call.

# How do I stop it?

What if a process is trapping SIGINT?

- ▶ `^\`— sends SIGQUIT
- ▶ `kill -3 PID / kill -QUIT PID`— sends SIGQUIT
- ▶ `kill -9 PID` — sends SIGKILL
  - ▶ SIGKILL can't be blocked or caught

`kill -9` all the time?

---

<sup>3</sup>Usually

# How do I stop it?

What if a process is trapping SIGINT?

- ▶ `^\\`— sends SIGQUIT
- ▶ `kill -3 PID / kill -QUIT PID`— sends SIGQUIT
- ▶ `kill -9 PID` — sends SIGKILL
  - ▶ SIGKILL can't be blocked or caught

`kill -9` all the time?

- ▶ Better<sup>3</sup> to give processes a chance to clean up.

---

<sup>3</sup>Usually

# SIGCHLD

`Wait()` will block until a child ends.

Not useful if you want to do other work while the child is running.

- ▶ You could use:

```
result = waitpid(pid, &status, W_NOHANG)
```

To check the status of a particular child.

- ▶ What if there are lots of children? Need to check each one?
- ▶ SIGCHLD is sent to the parent process when something happens to one of its children.
- ▶ You can set a handler to act when SIGCHLD arrives.



# SIGCHLD

See `child.c`

# Notes

- ▶ Ideally your signal handlers should not:
  - ▶ take a long time to execute
  - ▶ acquire locks
- ▶ `sigwait()` may be useful if your main program needs to wait for a signal.
- ▶ Can't we just do ... and zombies will never be created?
  - ▶ The documentation says that techniques vary across systems.
- ▶ "Core dumps?" — in ye olden days segfaults (and some other signals) would cause a copy of memory (core) to be put into a file:
  - ▶ These were useful for
    - ▶ debugging
    - ▶ using all your quota
  - ▶ Most modern linux systems have core dumps disabled.

## Other signals

- ▶ SIGSEGV — surprise segfault
- ▶ SIGWNCH — your terminal window changed size