

## CSSE2310 — 8.1

### Threads and Synchronization

# Main challenges

- ▶ Process management functions are weird.
- ▶ Debugging at a distance
- ▶ Debugging multiple processes simultaneously
- ▶ Being disciplined in the way you code and test.

# Threads

## Not for Ass3

- ▶ Threads and synchronization is for Ass4.
- ▶ Do not use this material in Ass3!

# Working in a process

To run code, we need:

- ▶ A process container
  - ▶ Address space
  - ▶ open files
  - ▶ ...
- ▶ Code / Instructions
  - ▶ “text segment”
- ▶ Global variables / constants
- ▶ Heap
  - ▶ Dynamic storage
- ▶ Stack
  - ▶ Function local variables
- ▶ CPU and its registers

## Multiple workers?

Suppose we want to have multiple workers doing things? `fork()`?

- ▶ Open files are shared (see Week 7.1).
- ▶ Everything else is separate.
- ▶ Needed if you want to have the new worker running a different program.

Different processes can cooperate on tasks (especially since they can start off with the same knowledge).

- ▶ But communicating results is more complicated.
  - ▶ Pipes?
  - ▶ “Network connection”? Similar complexity to pipes.
  - ▶ Shared memory?
  - ▶ ...
- ▶ But processes are safe from each other

# Threads?

A thread is a worker in a process<sup>1</sup>.

So every process has at least one thread (even if it doesn't explicitly use a thread library).

- ▶ Likely you will only interact with these via libraries.
- ▶ Different systems may have distinctions between “Threads / Light weight processes / ...”
  - ▶ We will avoid talking about these
- ▶ Threads can be classified as “native” or “green”
  - ▶ We'll get to this later

---

<sup>1</sup>A thread of execution through the program

# Multiple threads

If we want multiple threads operating in the same **process**, we would want them to be able to share information easily.

Item	Shared?
Process	Yes
Code	Yes
Global variables	Yes
Heap	Yes
Stack	No
CPU + registers	No





# Threads and memory

Any thread running in a process can interact with any variable used by any other thread if they know where it is (ie they have a pointer to it).

The kernel won't stop this because, if a page is valid for one thread, it is in the page table and so valid for all of the other threads.

# Thread implementations

Two main approaches here:

- ▶ “native” or “kernel” threads<sup>2</sup>
  - ▶ Threads are known by the kernel.
  - ▶ Threads can be scheduled and run independently of each other
    - ▶ Scheduling is out of scope for 2310, but gang scheduling is an interesting example.
- ▶ “green” threads
  - ▶ Kernel just sees a single threaded process
  - ▶ User space library switches saves registers and modifies program counter to switch between activities
  - ▶ Not actually using multiple cpu cores
  - ▶ A kernel call in one means you can't switch until you get back to user mode
    - ▶ Blocking I/O?
  - ▶ eg Python threading

---

<sup>2</sup>We are assuming these for 2310

# Why threads?

Why do we teach threads? Are they actually needed?

Threads are:

1. Things you should know about<sup>3</sup>
2. A way around blocking I/O
  - ▶ I need to listen to 5 network connections<sup>4</sup> and I don't know what order they will return in.
3. A way to use more than one core worth of CPU power.

---

<sup>3</sup>ie a well trained programmer should know about them.

<sup>4</sup>Blocking reads

## 2 Blocking I/O

Situation:

Need to interact with multiple I/O FILE\*/fds

```
while (!done) {  
    read from A and process  
    read from B and process  
    read from C and process  
}
```

Problem: If input is available from C but not A, we won't get there.

## 2 Blocking I/O

Threading approach:

```
... interact(FILE* src) {  
    while (... ) {  
        read from src  
        process input  
    }  
}  
  
...  
    run interact(A) in thread    // not actual syntax  
    run interact(B) in thread  
    run interact(C) in thread
```

Each thread blocks on its own input source and doesn't interact with others.

## Non-threaded alternative?

Non-blocking I/O calls exist in C, so you could do this:

```
while (!done) {  
    if (have_input_on(A)) {  
        read from A and process  
    }  
    if (have_input_on(B)) {  
        read from A and process  
    }  
    if (have_input_on(C)) {  
        read from A and process  
    }  
}
```

Problem: **busy waiting**. This loop may run hundreds(?) of times before any input arrives.

# Event driven programming

```
while (!done) {  
    sleep_until_input_available_on(A, B, C)  
    process input  
}
```

These sorts of calls exist and deal with the busy wait problem. One core may be plenty — since most of the time is spent waiting. So why aren't we using them?

- ▶ You need to work with `fds` and `read()`. No `FILE*` niceties.
- ▶ It is trickier to keep track of where you are up to with multiple tasks (instead of each task having its own thread context).

## More power

In applications which do not do much I/O (eg scientific computing), non-blocking operations are not going to help. There you either need multiple processes or multiple threads to access more cores.



# pthread

C has no standard threading library.



- ▶ We are using the **POSIX-thread** API (pthread).
- ▶ Implementations exist for most OS (including windows)
- ▶ pthread wrap OS threads calls.

From now on, any discussion of “threads” means “pthread threads”.

# pthread

- ▶ Threads have no parent child relationships.
- ▶ Any thread can “join” (ie wait on) any other thread in the same process.
- ▶ Threads can't interact with threads in other processes.
- ▶ Any thread calling `exit()` will end the whole process including all other threads.
- ▶ `fork()`?
  - ▶ Yes that is an interesting question.
- ▶ Signal handlers?
  - ▶ Yes ... later

## thread functions

- ▶ Each thread is created to run a function.
- ▶ When that function ends / returns, so does the thread<sup>5</sup>
- ▶ The function needs to have the correct signature.

```
void* foo(void*) {  
    // do thread things  
    return (void*)0;  
}
```


So the function can

- ▶ Accept any parameter (as long as you can express it as a `void*`)
- ▶ Send back anything to whoever `join()`s on it (as long as it's a `void*`)

---

<sup>5</sup>As far as you know.

See thread1.c

Note: to compile a program using pthreads you need to add  
-pthread to the command line<sup>6</sup> 

```
joel@sage:~/csse2310_2020_1/lecture/code$ cat thread1.c
#include <pthread.h>
#include <stdio.h>
#include <unistd.h> // for sleep()
```

```
void* hello(void* v) {
    char* s = (char*)v;
    printf("Hello %s\n", s);
    return 0;
}
```

```
int main(int argc, char** argv) {
    pthread_t tid;
    pthread_create(&tid, 0, hello, "Larry");
    pthread_create(&tid, 0, hello, "Curly");
    pthread_create(&tid, 0, hello, "Moe");
    sleep(2);
    return 0;
}
```

```
joel@sage:~/csse2310_2020_1/lecture/code$ gcc -pthread thread1.c
joel@sage:~/csse2310_2020_1/lecture/code$ ./a.out
Hello Larry
Hello Curly
Hello Moe
joel@sage:~/csse2310_2020_1/lecture/code$ ./a.out
Hello Curly
Hello Larry
Hello Moe
```

# Hello

man pthread\_create



```
int pthread_create(pthread_t* thread,
    const pthread_attr_t* attr,
    void* (*start_routine)(void*),
    void* arg);
```

- ▶ the id of the new thread is stored at the first argument
  - ▶ if the call is successful — see return value.
- ▶ we'll ignore the second argument for now
- ▶ the function we want to call
- ▶ the argument we want to pass it



# Hello

What is the call to `sleep()` doing there?

See `thread2.c`

We can end an individual thread using `pthread_exit()`;

See `thread3.c`

```
#include <stdio.h>
#include <unistd.h> // for sleep()
```

```
void* hello(void* v) {
    char* s = (char*)v;
    printf("Hello %s\n", s);
    return 0;
}
```

```
int main(int argc, char** argv) {
    pthread_t tid;
    pthread_create(&tid, 0, hello, "Larry");
    pthread_create(&tid, 0, hello, "Curly");
    pthread_create(&tid, 0, hello, "Moe");
    return 0;
}
```

```
joel@sage:~/csse2310_2020_1/lecture/code$ gcc -pthread thread2.c
joel@sage:~/csse2310_2020_1/lecture/code$ ./a.out
joel@sage:~/csse2310_2020_1/lecture/code$ ./a.out
joel@sage:~/csse2310_2020_1/lecture/code$ ./a.out
joel@sage:~/csse2310_2020_1/lecture/code$
```

```
int main(int argc, char** argv) {
    pthread_t tid;
    pthread_create(&tid, 0, hello, "Larry");
    pthread_create(&tid, 0, hello, "Curly");
    pthread_create(&tid, 0, hello, "Moe");
    pthread_exit((void*)0);
}
```

```
joel@sage:~/csse2310_2020_1/lecture/code$ gcc -pthread thread3.c
joel@sage:~/csse2310_2020_1/lecture/code$ gcc -pthread -Wall -pedantic -std=gnu99 thread3.c
joel@sage:~/csse2310_2020_1/lecture/code$ ./a.out
Hello Larry
Hello Curly
Hello Moe
```

# Passing values into threads

So far we've passed in strings (char\*).

► char\* → void\* → char\*

What about int?

See thread4.c

Why bother with malloc()ing?

See thread5.c

This is an example of a “race condition”.



```
#include <stdio.h>
#include <unistd.h> // for sleep()
#include <stdlib.h> // malloc and free
```

```
void* hello(void* v) {
    int value = *(int*)v;
    free(v);
    printf("Hello %d\n", value);
    return (void*)0;
}
```

```
int main(int argc, char** argv) {
    pthread_t tid;
    for (int i=0; i<5; ++i) {
        int* val = (int*) malloc(sizeof(int));
        *val = i;
        pthread_create(&tid, 0, hello, val);
    }
    pthread_exit((void*)0);
}
```

```
joel@sage
Hello 2
Hello 2
Hello 3
Hello 1
Hello 4
Hello 0
joel@sage
Hello 2
Hello 1
Hello 0
Hello 4
Hello 3
joel@sage
Hello 1
Hello 2
Hello 0
Hello 4
Hello 2
Hello 3
Hello 4
```

```
joel@sage
Hello 2
Hello 2
Hello 3
Hello 4
Hello 4
joel@sage
Hello 1
Hello 2
Hello 4
Hello 4
Hello 4
Hello 4
```

The Unix

```
void* hello(void* v) {
    int value = *(int*)v;
    printf("Hello %d\n", value);
    return (void*)0;
}

int main(int argc, char** argv) {
    pthread_t tid;
    for (int i=0; i<5; ++i) {
        int val = i;
        pthread_create(&tid, 0, hello, &val);
    }
    pthread_exit((void*)0);
}
```