

CSSE2310 — 8.2

Threads and Synchronization continued

Passing values into threads


So far we've passed in strings (`char*`).

▶ `char* → void* → char*`

What about `int`?

See `thread4.c`

Why bother with `malloc()`ing?

See `thread5.c` 

This is an example of a “**race condition**”.

Abusing pointers

ints fit in void* right?

```
void* hello(void* v) {  
    int value = (int)v;  
    printf("Hello %d\n", value);  
    return (void*)0;  
}  
  
int main(int argc, char** argv) {  
    pthread_t tid;  
    for (int i=0; i<5; ++i) {  
        pthread_create(&tid, 0, hello, i); // implicit int -> void*  
    }  
    pthread_exit((void*)0);  
}
```

What about thread6.c?

```
thread6.c:7:17: warning: cast from pointer to integer of different size [-Wpointer-to-int-cast]  
    int value = (int)v;  
                  ^
```

- ▶ We could pass long instead to get around the size issue.
- ▶ The real problem¹ here is that it works
 - ▶ ...in this case.
 - ▶ it encourages people to ignore warnings

```
int main(int argc, char** argv) {  
    printf("%d %d %d %d %d\n", sizeof(short), sizeof(int), sizeof(long), sizeof(long long), sizeof(void*));  
    return 0;  
}  
  
joel@sage:/tmp$ ./a.out  
2 4 8 8 8
```

¹from a teaching point of view

Multiple parameters

Suppose we have a function which we would like to run in a thread:

```
void do_things(char** items, char* s, int limits);
```

It doesn't match the required signature.

Multiple parameters

Declare a struct type to hold all the values:

```
struct Params {  
    char** items;  
    char* s;  
    int limits;  
};
```

Add a wrapper function

```
void* do_things_wrapper(void* v) {  
    struct Params* p = (struct Params*)v;  
    do_things(p->items, p->s, p->limits);  
    free(v);    // can't free earlier without copying the struct  
}
```

Elsewhere:

```
struct Params* p = malloc(sizeof(struct Params));  
p->items = items1;  
p->s = "target";  
p->limits = 10;  
pthread_create(&tid, 0, do_things_wrapper, p);
```

What's a thread id?

- ▶ `<pthread.h>` declares a `pthread_t` type.
- ▶ It is an *opaque* type
 - ▶ Makes printing it in debugging awkward
 - ▶ (On `moss` it turns out to be `unsigned long` but you can't rely on that)

Where does thread return go?

```
int main(int argc, char** argv) {
    pthread_t tid;
    int* p = malloc(sizeof(int));
    *p = 4;
    pthread_create(&tid, 0, do_cube, p);
    // Now we wait for the thread to finish
    // We need somewhere to store the return value
    void* res;
    pthread_join(tid, &res); // &res is void**
    printf("Thread returned %d\n", *(int*)res);
    return 0;
}
joel@sage:~/csse2310_2020_1/lecture/code$ gcc -pthread thread_calc.c
joel@sage:~/csse2310_2020_1/lecture/code$ ./a.out
Thread returned 64
```

See thread_calc.c

Note:

- ▶ Since the thread function returns void*, to allow pthread_join to modify a variable, you need to pass in void**.
- ▶ Most pthread functions² return error codes, which you should check.
- ▶ pthread_exit(V) ends the thread and sets its result to V;



```
// In this limited case we could have used the same pointer
// for in and out
```

```
void* alt_cube(void* v) {
    int* p = (int*)v;
    int value = *p;
    *p = value * value * value;
    return (void*)0;
}
```

```
void* do_cube(void* v) {
    int val = *(int*)v;
    free(v);
    int* result = malloc(sizeof(int));
    *result = val * val * val;
    return (void*)result;
}
```

²in particular pthread_create() and pthread_join()

Zombie threads?

- ▶ Yes, the doco indicates these exist.
- ▶ `pthread_join()` deals with them in a similar way to `wait()` reaping zombie processes.
- ▶ Threads don't terminate independently due to signals
 - ▶ Unhandled signals will take out the whole process.
- ▶ `pthread_detach(tid)` tells the system to clean up zombie threads automatically.
 - ▶ You can use the second argument of `pthread_create` to start a thread detached.

Killing threads?

Can you “kill” threads?

- ▶ (as previously noted) not using signals.
- ▶ `pthread_cancel(tid)` exists.

In general, it is difficult to do safely.

A “safer” approach is to use a shared variable that the thread checks.

```
void* thread_fn(void* v) {  
    bool* stop=...           // get this some how  
    // other things  
    while (!*stop) {  
        // do some work  
    }  
}
```

So, yes you can but you shouldn't.

pthread_self()


```
pthread_t pthread_self(void);
```

Returns the thread id of the current thread.

Racing

```
joel@sage:~/csse2310_2020_1/lecture/code$ gcc -pthread race1.c
joel@sage:~/csse2310_2020_1/lecture/code$ ./a.out
Expected 1200000 got 220736
Lost 979264 updates
```

See race1.c

Problem: value++ is not  **atomic**³

```
#define LOOPS 200000
#define PCOUNT 6

void* fn(void* v) {
    int* p=(int*)v;
    for (int i=0; i < LOOPS; ++i) {
        (*p)++;
    }
    return 0;
}

int main(int argc, char** argv) {
    int total = 0;
    pthread_t tids[PCOUNT];
    for (int i = 0; i < 6; ++i) {
        pthread_create(&(tids[i]), 0, fn, &total);
    }
    for (int i = 0; i < 6; ++i) {
        void* v;
        pthread_join(tids[i], &v);
    }
    printf("Expected %d got %d\n", LOOPS*PCOUNT, total);
    printf("Lost %d updates\n", LOOPS*PCOUNT-total);
    return 0;
}
```

³ie indivisible (greek atomos)

Racing

Let's use a lock.

See `race2.c`

Hmm... Our lock code is not atomic either.



Running `race3.c` shows that locks can be made to work.

Thread coordination

Problems

The `race2` implementation fails to do two desirable things:

1. Ensure mutual exclusion



- ▶ Much simpler with hardware assistance.
- ▶ Normally access this via library calls.

2. Avoid busy waiting

Note: For a mutex algorithm that doesn't rely on⁴ hardware assistance see Peterson's algorithm.

⁴much

Semaphores

An opaque type which represents an integer value.

Two atomic operations :

- ▶ `sem_wait()`
 - ▶ If the value is > 0 decrement it by 1.
 - ▶ If the value $== 0$, stop the process until value > 0 , then attempt to decrement.
- ▶ `sem_post()`
 - ▶ Increment the value.

If the p threads are waiting on the semaphore and a `post()` occurs, only one of the threads unblocks, the others stay blocked.

Does everyone get a turn eventually?

- ▶ “Starvation” is beyond the scope of this course



Semaphores in action

See `race3.c`

- ▶ `sem_init()` — set the initial value of the semaphore
- ▶ `sem_wait()` — decrement the semaphore
- ▶ `sem_post()` — increment the semaphore
- ▶ `sem_destroy()` --- clean up

Note: Always pass pointers to semaphores, do not copy the value itself.