

CSSE2310 — 9.1

Threads and Synchronization continued

Semaphores in action

See `race3.c`

- ▶ `sem_init()` — set the initial value of the semaphore
- ▶ `sem_wait()` — decrement the semaphore
- ▶ `sem_post()` — increment the semaphore
- ▶ `sem_destroy()` --- clean up

Note: Always pass pointers to semaphores, do not copy the value itself.

Mutual exclusion

Task 1: mutex

- ▶ `sem_init()`: set value to 1
- ▶ `sem_wait()`: acquire the lock / mutex. Only one thread can succeed until post
- ▶ `sem_post()`: release the lock

Provided that all paths into the **critical section(s)**¹² require waiting and all paths out post, this will ensure mutual exclusion.

¹No more than one thread should be in a critical section at a time.

²Different parts of a program could belong to the same critical section.

Waiting

Task 2: Non-busy waiting



- ▶ `sem_init()`: set value to 0
- ▶ `sem_wait()`: block until semaphore is available.
- ▶ `sem_post()`: let other thread know “it” happend.



Other things to do with semaphores

Limit maximum threads active (not as common)

- ▶ `sem_init()`: set value to N
- ▶ at most N threads can pass (wait) until one or more threads leave (post).

Producer and consumer tasks

- ▶ `sem_init()`: set value to 0
- ▶ Each time the “producer” adds a job to the queue, post
- ▶ Consumer threads all wait on the semaphore.

Notes:

- ▶ in this case some threads only wait and other threads only post
- ▶ You still need a separate mutex to control accessing the queue.

Volatility

Consider the following code:

```
total=0;
if (*a > 0) {
    total++;
}
if (*a > 0) {
    total++;
}
if (*a > 0) {
    total++;
}
// total is either 0 or 3?
```

Volatility

Or:

```
total = 0;
ref = *a;
total += *a;
total += *a;
total += *a;
total += *a;
// total is 4 * ref?
```

The compiler won't see any reason that it can't optimise the calculation (and most of the time it would be correct).

- ▶ What if another thread also knows pointer *a*?
- ▶ ... and modifies it

Volatility

The `volatile` keyword warns the compiler that the value of the variable may change in ways (and at times) when the compiler won't predict.

- ▶ ie. Don't get too clever with this one.

eg:

```
volatile int* p;  
           // p is a pointer to a volatile int  
volatile const int x=6;  
           // yes this is legal  
const int* const p=&global;  
           // not volatile but still legal
```


Not volatile

Use `volatile` when any variable may be modified by one thread and read in others.

```
void f (...) {  
    int a = ...;  // not volatile  
}
```

Even if multiple threads call `f()` at the same time, they will each see a different (local) variable.

thread safety

An operation is “thread safe”³ if multiple threads can have active calls to the function at the same time.⁴

Things to look for:

- ▶ A value could be modified by one thread while another is using it.
 - ▶ This includes freeing and mallocing, removing entries from lists, ...
 - ▶ Can be tricky to spot (eg `i++`) if you are used to thinking of them as atomic

³You may also see this called “reentrant”

⁴Could also consider recursion or nested signal handlers here

thread safety

Look for calls to non-threadsafe functions.

- ▶ How to tell? Look at the (up to date) doco.
- ▶ eg: `rand_r`
- ▶ `_r` normally indicates “reentrant”
- ▶ `man 3 rand_r`
- ▶ On moss, says `rand()` and `rand_r()` are both thread safe.
- ▶ (and that `rand_r` is deprecated).

Some functions make use of hidden static state which might not be obvious.

Non-obvious non-atomics

Suppose many threads execute the following at the same time:


```
*p = xn ;
```

Thread 1 stores x_1 , ... Thread n stores x_n .

What happens?

- ▶ There is a race condition where it is not known which thread will write last.
- ▶ Will $*p$ store one of $\{x_1, \dots, x_n\}$?
 - ▶ It depends on the size of x_i vs the size of a processor word.
 - ▶ eg long double on moss.
 - ▶ long on 32bit systems?

Dining philosophers



Deadlock

Thread 1:



```
wait(l1);  
wait(l2);  
    // do something  
post(l1);  
post(l2);
```

Thread 2:

```
wait(l2);  
wait(l1);  
    // do something  
post(l2);  
post(l1);
```

- ▶ For simple cases: Everyone should request resources in the same order.
- ▶ More complex cases: beyond the scope of this course.

Dining philosophers

Potential problems:

- ▶ Deadlock
- ▶ Livelock
- ▶ Starvation⁵

⁵Again beyond the scope of the course

Misc pthreads


fork() and threads

According to the documentation, calling `fork()` in a multithreaded program only duplicates the thread which called `fork()`.

See `forknthread.c` as an example.

You must also be careful of locks etc.

semaphores

- ▶ `sem_trywait()`
 - ▶ Either lock immediately or error
- ▶ `sem_timedwait()`
 - ▶ Error if lock can't be acquired before timeout 
 - ▶ Note: this function takes an **absolute time** not a delta.
- ▶ semaphores between processes? (Not in 2310).
 - ▶ `sem_create()` / `sem_destroy()` in shared memory
 - ▶ `sem_open()` / `sem_unlink()` named semaphores

mutex and condition vars

threads also specifies:

- ▶ `pthread_mutex_t`
 - ▶ Only mutual exclusion
 - ▶ Can only unlock from the thread which locked
- ▶ `pthread_cond_t`
 - ▶ For waiting for something
 - ▶ Each condition var is linked to a mutex