

CSSE2310/7231 — A.1

Layers continued & Programming

Note

Remember there is a distinction between *bandwidth* and *latency*.

- ▶ Bandwidth = how much information you can send per unit time. (bits/s)
- ▶ Latency = How long it takes for something to start (Time)

Layers

- 5 Application
- 4 Transport
- 3 Network
- 2 Data link
- 1 Physical

Where do addresses come from?

- ▶ ports — low ports aside, you can say which port you want to use when you set up a network connection.
- ▶ IP —
 - ▶ internal configuration
 - ▶ The administrator of the device chooses one.
 - ▶ If they choose badly it might not work. (more later)
 - ▶ external configuration
 - ▶ On startup the device can ask another device what address it should use
 - ▶ DHCP (more later)
- ▶ MAC
 - ▶ interfaces will have a default address set in hardware.
 - ▶ The address that is used could be the default or could be changed by software.

IPv4 has 32bit addresses. In theory that's 4,294,967,296 addresses.

- ▶ N.A.T.¹ has probably helped reduce demand a bit.
- ▶ Addresses weren't allocated very efficiently.
 - ▶ US military has at least 134,217,728 reserved².
 - ▶ Apple, Ford and Daimler have 16,777,216 each.

For a while now, big blocks of addresses were handed out to regional "registries" and organisations which divided them up into smaller blocks to give to smaller organisations. . . . all the way down to your ISP giving you an address.

¹later

²That's a decrease from what they had

We need more addresses

- ▶ The world ran out³ of IPv4 addresses in 2011.
- ▶ The solution the IETF had proposed was IPv6 with 128bit addresses.

That's

... 340, 282, 366, 920, 938, 463, 463, 374, 607, 431, 768, 211, 456
addresses.

Which is a lot.

³At the top level anyway

Why aren't we using IPv6 in the teaching and assessment in this course?

- ▶ It is possible to write code which is IP version agnostic.
 - ▶ So you wouldn't really learn anything different (at this level)
- ▶ Do you want 39 digit numbers in your exams?
 - ▶ Yes I know about the :: contraction convention

client / server

Client ↔ Server

Server — a **process** which waits for requests from clients.

- ▶ eg: A web server waits on port 80 for browsers to connect and request pages.
- ▶ `sshd` waits for connections on port 22.

Client — a **process** which submits requests to a server.

- ▶ A web client connects **to** a server on port 80 and requests pages.
- ▶ `ssh` or `putty` connects to the server on port 22.

Client \leftrightarrow Server

Unfortunate terminology:

- ▶ “server” refers to a process running on a machine.
 - ▶ Note: Server on port x , not port x on the server.
- ▶ “server” is also often used to refer to the hardware that the server process runs on.
 - ▶ We run servers on your servers . . .

TCP Connections

- ▶ The “client / server” distinction applies to single connections.
- ▶ For any TCP connection there is always a client and a server.
- ▶ An application acting in a peer-to-peer mode, may have some connections for which it was the server and others for which it was the client.
- ▶ Once a connection has been established, there is no difference between what a client can do and what a server can do.

ncat / nc

You want to be able to debug your network code without needing both your client and your server working. **Netcat to the rescue.**

To start a server (it's `-e11`):

```
$ nc -4 -l 43210
```

Using IPv4 listen on port 43210.

To connect to your server

```
$ nc -4 localhost 43210
```

Using IPv4 connect to port 43210 on a computer called localhost.

Notes

- ▶ Connections are bi-directional (you can type into either and it gets sent to the other end).
- ▶ Ports are machine-wide so you will need to pick a number that noone else is using.
- ▶ On `localhost` you can also use the machine name `localhost4` which only has a IPv4 address.
- ▶ `localhost` is defined on most systems to give **an** IP address of the machine you are on.
 - ▶ Remember, machines can have multiple names and multiple IPs

Warning:

- ▶ There are multiple versions of netcat out there.
- ▶ We will be using the version installed on moss. (In debian this is in the `netcat-openbsd` package.
- ▶ You may need to check the doco for the version on your system.

Client steps

1. Find out the address of the machine you wish to connect to.
2. Make a socket
 - ▶ A socket is an abstraction of a network endpoint.
 - ▶ There are different types of socket, we'll be using `SOCK_STREAM`
 - ▶ Sockets are file-descriptors
3. `connect()` to the server.
4. Wrap socket descriptor for nicer IO
 - ▶ You should `dup()` the descriptor before calling `fdopen()`

net1.c

See net1.c

So struct addrinfo

- ▶ contains a struct sockaddr*
- ▶ which is actually a struct sockaddr_in*
 - ▶ which contains a struct in_addr
 - ▶ which contains a in_addr_t which is an unsigned integer of some sort.

Why?

- ▶ The levels of indirection may seem excessive, but most of the time you won't be dealing with directly with all that detail.
- ▶ This allows different types of sockets to be used.
- ▶ Network byte order may be different to your machine's natural order.

See `net2.c`

Server

See `net3.c`

Server steps

1. Make a socket
2. (Optional) set parameters
3. `bind()` the socket to a port.
 - ▶ Otherwise it becomes:
 - ▶ “I have a server can you guess where?”
4. Set the socket to `listen()` for connections.
5. Call `accept()` to get allow a connection
 - ▶ Use the new fd to interact with the client

See `net3b.c` for options.

See `net4.c`

This would not be normal for a server

- ▶ normally you want people to connect to a known port
- ▶ we are doing it to avoid port collisions during the assignment.

ntohs?

`ntohs` = **n**etwork **t**o **h**ost **s**hort

Converts a 16 bit value from network representation to the machines normal ordering.

Others:

- ▶ `ntohl`
- ▶ `htons`
- ▶ `htonl`

Multi-programming

`accept()` is a blocking call. If you want to be able to work on a connection **and** wait for new connections you need extra workers.
Possibilities:

- ▶ `fork()` and let the new child handle the new connection.
- ▶ create a `pthread` to handle the new connection.
- ▶ use a non-blocking call to check for IO and connections (**not in this course**).