

CSSE2310

...

Concurrency and Threads

Week 11, Sem 1, 2020

Week 11

- Assignment 4 is being prepared for release
 - Involves threads, mutexes and semaphores -> this tute!
 - Involved network programming
- Useful resources
 - Lectures (concurrency and network programming)
 - Modularity PDFs on Blackboard
 - Piazza (see if your question has been asked already)
- Today's content
 - Thread diagrams and concurrency material
 - Running through an example concurrent program and making it thread-safe

Threads: A Refresher

- You can create a thread using `pthread_create`:

```
int pthread_create(pthread_t* tid, const pthread_attr_t* attr,  
                  void* (*start)(void*), void* arg)
```

- First argument is a pointer to store the ID of the created thread
- Second argument is a pointer to a struct containing thread attribute settings (usually NULL)
- Third argument is a pointer to the function which you want the thread to start executing at
- Last argument is the `void*` to the data you want to pass as parameters to the thread function

Exercise: Thread Diagram Question 1

- What diagram is produced by the following code? What is the execution time?

```
int main(int argc, char** argv) {  
  
    pthread_t tId[2];  
    pthread_create(&tId[0], NULL, thread1,  
                  NULL);  
    sleep(1);  
    pthread_create(&tId[1], NULL, thread2,  
                  (void*) tId[0]);  
    sleep(1);  
    pthread_join(tId[1], NULL);  
    printf("All done\n");  
    return 0;  
}
```

```
void* thread1(void* arg) {  
    sleep(1);  
    printf("Thread 1 ending\n");  
    return NULL;  
}  
  
void* thread2(void* arg) {  
    sleep(1);  
    pthread_join((pthread_t) arg, NULL);  
    printf("Thread 2 ending\n");  
    return NULL;  
}
```

Exercise: Thread Diagram Question 1 Solution


```
int main() {  
    pthread_t tId[2];  
    pthread_create(&tId[0], NULL,  
thread1,  
        NULL);  
    sleep(1);  
    pthread_create(&tId[1], NULL,  
thread2,  
        (void*) tId[0]);  
    sleep(1);  
    pthread_join(tId[1], NULL);  
    printf("All done\n");  
    return 0;  
}
```

```
void* thread1(void* arg) {  
    sleep(1);  
    printf("Thread 1 ending\n");  
    return NULL;  
}  
  
void* thread2(void* arg) {  
    sleep(1);  
    pthread_join((pthread_t) arg, NULL);  
    printf("Thread 2 ending\n");  
    return NULL;  
}
```

The pink sleep and first main sleep will run in parallel, so is expected to take 1 second in total.



The purple sleep and the second main sleep will also run in parallel, so is expected to take 1 second in total.

Therefore we expect the program to take 2s.




Exercise: Thread Diagram Question 2

- Which of the following are possible with threads?



Exercise: Thread Diagram Question 2 Solution


- Which of the following are possible with threads?



This is possible with threads, however we need a global variable for A to know C's thread id so it can join with it.

Exercise: Thread Diagram Question 2 Solution


- Which of the following are possible with threads?



This is also possible with threads, see the thread diagram example.

Exercise: Thread Diagram Question 2 Solution


- Which of the following are possible with threads?




This is also possible with threads, this was covered in the concurrency lectures and uses `pthread_exit`.

Concurrency

- Two actors* are concurrent if they could run at the same time
 - This doesn't necessarily mean that they run in parallel
 - It just means that there isn't a guaranteed ordering in how they run
- Example: Concurrent actors which aren't necessarily run in parallel



2 concurrent actors using 2 CPUs



2 concurrent actors using 1 CPU

*processes, threads (in other environments these could be tasks etc.)

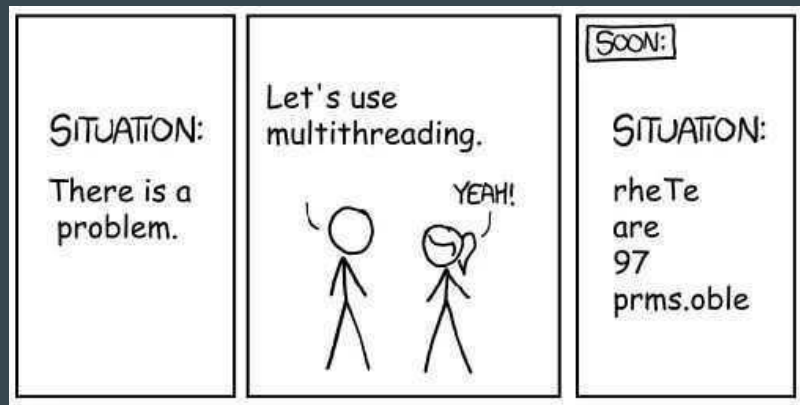
Concurrency: Pros and Cons

Pros

- Can be faster
- More efficient CPU utilisation
- Lets you do more than one thing at a time. For example:
 - More than one visitor can view your website at once
 - More than one person can be logged into moss at once
 - An automated car can look at traffic lights while staying in its lane

Cons

- Race conditions
- Memory corruption... but worse and harder to debug than normal
- Non-deterministic bugs
- Debuggers hiding bugs



Concurrency: Race Conditions

- When A should happen before B, but the order gets confused



- Fixed by synchronising actors
 - Ensure that B doesn't start until A has finished
 - A must tell B when to start
 - Read / write from a pipe
 - Lock / unlock a mutex
 - Wait / post a semaphore
 - etc.

Concurrency: Memory Corruption

- One thread writes to memory while the other one tries to access it (read or write)
- This was avoided in assignment 3 due to **memory isolation**
 - You can't write to the same memory if that memory isn't shared
- Can be prevented by:
 - Using read-only memory (note that calling `free(3)` is a type of write)
 - Giving memory to another thread and no longer accessing it
 - Only accessing memory with a single thread at a time (via mutexes and semaphores)

Concurrency: Primitives

- Mutexes
 - Can only be “locked” by one thread at a time
 - If attached to a block of code, ensures only one thread runs that code at any one time (i.e. provides mutual exclusion)
- Semaphores
 - Based off the concept of a counter
 - Once the counter reaches zero, threads trying to decrement (wait) it are put to sleep
 - When a process increments the counter (post), one of the waiting processes will wake up

Concurrency: Deadlocks

- Deadlocks occur when different actors are waiting on each other in a **cycle**
 - A waiting for B, B waiting for C, C waiting for A
 - A waiting for B, B waiting for A
- Fix one: don't have cycles
 - Draw a diagram of which actor waits for which and ensure it is cycle free
 - Not always applicable, but reducing loops can simplify things
- Fix two: ensure at least one actor writes (and flushes) before they read
 - Draw some diagrams of communication over time, perform some tests
- When thinking about deadlocks, please remember
 - Reading is a form of waiting if there is something on the other end
 - Other processes, as well as threads, can be actors and can deadlock

Exercise: Making an I/O Program Thread Safe

- We've got a program that does I/O using multiple threads
- There are two "writing" threads:
 - One reads from stdin and sends the read data to the channel
 - One reads from a file and sends the read data to the channel
- And one "reading" thread:
 - Reads from the channel, and prints to stdout
- The channel in this case is a queue which is used as a means of communication between threads (i.e. it is essentially a fancy pipe - threads can write to it, and other threads can read from it)

Exercise: Making an I/O Program Thread Safe

- The code is available on Blackboard under Learning Resources -> Prac Slides
- The code contains:
 - A makefile (notice -g and -pthread)
 - main.c (logic for each of the threads)
 - queue.c/h (a FIFO queue implementation)
 - channel.c/h (wrappers around the queue to make a thread-safe channel)
- Running ./channel seems to work correctly
- Try to stress test it: `cat /dev/urandom | base64 | fold -w 60 | ./channel`
- It will break eventually if you leave it running
 - Memory corruption is occurring due to the channel not being thread-safe

Exercise: Making an I/O Program Thread Safe

- Step 1: Add a mutex to the channel
- Mutexes are built for this scenario. There is a piece of data (the queue) that you want to make sure no two threads write to at once. You fix this by wrapping each access to that data in a lock/unlock statement

read_channel:

```
pthread_mutex_lock(lock);  
output = read_queue();  
pthread_mutex_unlock(lock);
```

write_channel:

```
pthread_mutex_lock(lock);  
output = write_queue();  
pthread_mutex_unlock(lock);
```

Exercise: Making an I/O Program Thread Safe

- We should now be passing the stress test
- In a separate terminal, try:
 - `ps -L -o pid,uid,cmd,pcpu # Prints pid, uid, cmd (normal output) AND %CPU used for each thread`
- Why is one of the threads taking up 100% CPU?
- Which thread is it?

Exercise: Making an I/O Program Thread Safe

- The writer thread is causing the issue
 - It is doing what is known as “busy polling”
 - If nothing is available to be read, it will keep looping and burning CPU cycles
 - Our program currently calls `read_queue` more than 10 million times a second on `moss`
- We need a way to put this thread to sleep until information is ready to be output

Exercise: Making an I/O Program Thread Safe

- Enter semaphores
 - We use a semaphore in this case to keep track of how many elements are in the queue
 - The semaphore starts at 0
 - We increment it (post) each time we write to the channel
 - We decrement it (wait) each time we read from the channel

read_channel:

```
sem_wait(signal);  
  
pthread_mutex_lock(lock);  
output = read_queue();  
pthread_mutex_unlock(lock);
```

write_channel:

```
pthread_mutex_lock(lock);  
output = write_queue();  
pthread_mutex_unlock(lock);  
  
sem_post(signal);
```

Things to Remember

- No sleeping (or blocking) with a locked mutex
- If one thread is writing, no other thread is nearby
- Just because your code works once doesn't mean it will work every time